

LIN570: HW10 – maxent (100pts)

YOUR NAME (UW NetID)

Due date: 11pm on Dec 10, 2019 (Tuesday)

All the example files are under `~/dropbox/19-20/570/hw10/examples/`.

1. **Q1 (55 points):** Create a MaxEnt POS tagger, `maxent_tagger.sh`.

- The command line is: `maxent_tagger.sh train_file test_file rare_thres feat_thres output_dir`
- The `train_file` and `test_file` have the format (e.g., `test.word_pos`): $w_1/t_1 \ w_2/t_2 \ \dots \ w_n/t_n$
- `rare_thres` is an integer: any words (in the `train_file` and `test_file`) that appear LESS THAN `raw_thres` times in the `train_file` are treated as *rare words*, and features such as `pref=xx` and `suf=xx` should be used for rare words (see Table 1 in (Ratnaparkhi, 1996)).
- `feat_thres` is an integer: All the w_i features (i.e., `CurrentWord=xx` features), regardless of their frequency, should be kept. For all OTHER types of features, if a feature appears LESS THAN `feat_thres` in the `train_file`, that feature should be removed from the feature vectors.
- `output_dir` is a directory that stores the output files from the tagger. Your script should create the following files and store them under `output_dir`:
 - `train_voc` (e.g., `ex_train_voc`): the vocabulary that includes all the words appearing in `train_file`. The file has the format “word freq” where freq is the frequency of the word in the training data. The lines should be sorted by freq in descending order. For words with the same frequency, sort the lines alphabetically.
 - `init_feats` (e.g., `ex_init_feats`): features that occur in the `train_file`. It has the format `featName freq` and the lines are sorted by the frequency of the feature in the `train_file` in descending order. For features with the same frequency, sort the lines alphabetically.
 - `kept_feats` (e.g., `ex_kept_feats`): This is a subset of `init_feats`, and it includes the features that are kept after applying `feat_thres`.
 - `final_train.vectors.txt` (e.g., `ex_final_train.vectors.txt`): the feat vectors for the `train_file` in the Mallet text format. Only features in `kept_feats` should be kept in this file.
 - `final_test.vectors.txt`: the feat vectors for the `test_file` in the Mallet text format. The format is the same as `final_train.vectors.txt`.

- `final_train.vectors`: the binary format of the vectors in `final_train.vectors.txt`.
- `me_model`: the MaxEnt model (in binary format) which is produced by the MaxEnt trainer.
- `me_model.stdout` and `me_model.stderr`: the stdout (standard out) and stderr (standard error) produced by the MaxEnt trainer are redirected and saved to those files by running command such as `mallet train-classifier --trainer MaxEnt --input final_train.vectors --output-classifier me_model > me_model.stdout 2> me_model.stderr`. The training accuracy is displayed at the end of `me_model.stdout`.
- `sys_out`: the system output file when running the MaxEnt classifier with command such as `mallet classify-file --input final_test.vectors.txt --classifier me_model --output sys_out`.

Your script `maxent_tagger.sh` should do the following:

- (a) Create feature vectors for the training data and the test data. The vector files should be called `final_train.vectors.txt` and `final_test.vectors.txt`.
- (b) Run `mallet import-file` to convert the training vectors into binary format, and the binary file is called `final_train.vectors`.
- (c) Run `mallet train-classifier` to create a MaxEnt model `me_model` using `final_train.vectors`
- (d) Run `mallet classify-file` to get the result on the test data `final_test.vectors.txt`.
- (e) Calculate the test accuracy

For step 2-4, you should use Mallet commands. For Step 5, if you don't want to write code for it, you can use the `vectors2classify` command, which covers step 3-5. In that case, you need to convert `final_test.vectors.txt` to the binary format first.

For the first step, you need to write some code. Features are defined in Table 1 in (Ratnaparkhi, 1996) (see MaxEnt slides). The following is one way for implementing this step:

- (a) create `train_voc` from the `train_file`, and use the word frequency in `train_voc` and `rare_thres` to determine whether a word should be treated as a *rare word*. The feature vectors for rare words and non-rare words are different.
- (b) Form feature vectors for the words in `train_file`, and store the features and frequencies in the training data in `init_feats`.
- (c) Create `kept_feats` by using `feat_thres` to filter out low frequency features in `init_feats`. Note that w_i features are NOT subject to filtering with `feat_thres` and every w_i feature in `init_feats` should be kept in `kept_feats`.
- (d) Go through the feature vector file for `train_file` and remove all the features that are not in `kept_feats`.
- (e) Create feature vectors for `test_file`, and use only the features in `kept_feats`. If a word in the `test_file` appears LESS THAN `rare_thres` times (or does not appear at all) in the `training_file`, the word should be treated as a *rare word* even if it appears many times in the `test_file`.
- (f) For the feature vector files, replace all the occurrences of “,” with “**comma**” as Mallet treats “,” as a separator.

2. **Q2 (20 points):** Run `maxent_tagger.sh` with `wsj_sec0.word_pos` as `train_file`, `test.word_pos` as `test_file`, and the thresholds as specified in Table 1:

- *training accuracy* is the accuracy of the tagger on the `train_file`
- *test accuracy* is the accuracy of the tagger on the `test_file`
- *# of feats* is the number of features in the `train_file` before applying `feat_thres`
- *# of kept feats* is the number of features in the `train_file` after applying `feat_thres`
- *running time* is the CPU time (in minutes) of running `maxent_tagger.sh`.

Table 1: Tagging accuracy with different thresholds

Expt id	rare thres	feat thres	training accuracy	test accuracy	# of feats	# of kept feats	running time
1_1	1	1					
1_3	1	3					
2_3	2	3					
3_5	3	5					
5_10	5	10					

Please do the following:

- Fill out Table 1.
- What conclusion can you draw from Table 1?
- Save the output files of `maxent_tagger.sh` to `res_id/`, where *id* is the experiment id in the first column (e.g., the files for the first experiment will be stored under `res_1_1`). Submit only the subdirs for the first row and the last row (i.e., `res_1_3` and `res_3_5`).

The submission should include:

- The `readme.[txt|pdf]` file that includes Table 1 and your answer to Q2.
- `hw.tar.gz` that includes `maxent_tagger.sh` and `res_1_3` and `res_3_5` created in Q2 (see the complete file list in `submit-file-list`).